

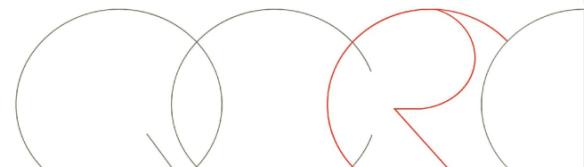


مختبر قطر لبحوث الحوسبة
Qatar Computing Research Institute

Graph Analytics

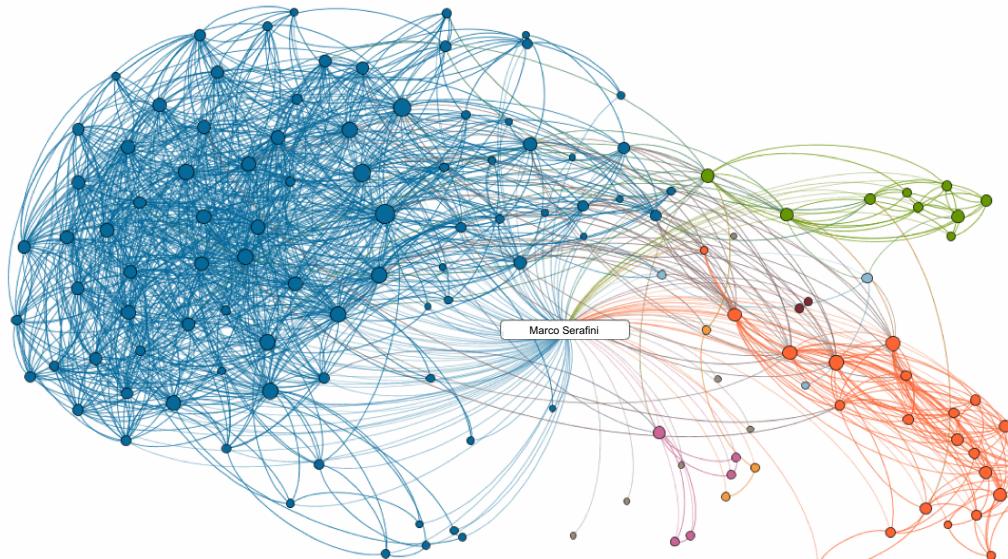
Graph Processing and Graph Mining

Marco Serafini, QCRI

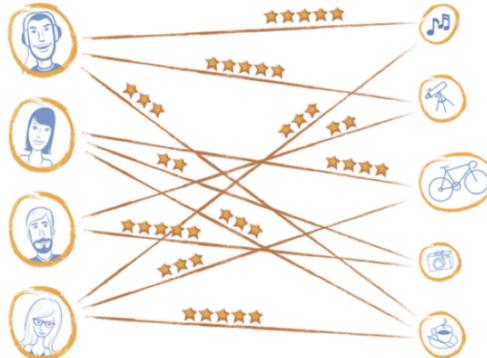


Graphs are Ubiquitous

LinkedIn Maps Georgos Siganos's Professional Network
as of November 11, 2013



©2013 LinkedIn - Get your network map at inmaps.linkedinlabs.com



Classes of Graph Systems

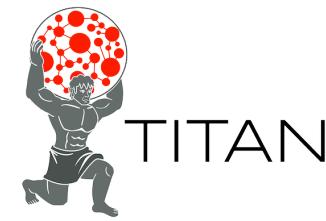
- **Graph Computation**

- Think like a vertex
- Linear algebra



- **Graph Search**

- Find instances of path expressions



- **Graph Mining**

- Mine patterns of interest and their matches



Graph Computation

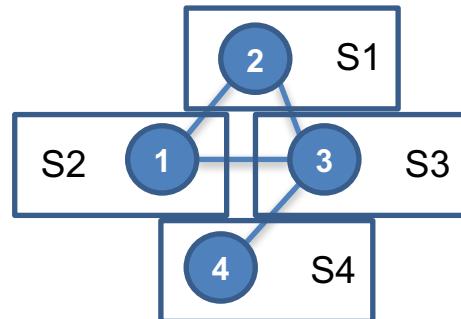


Examples of Graph Computations

- PageRank
 - Used to rank vertices (e.g., web pages) in a graph (e.g., of hyperlinks)
 - Each vertex is given a relevance score based on its centrality
- Shortest path
 - Used for example for routing
- Bipartite matching
 - Match vertices on different sizes of a bipartite graph
 - Use for example for scheduling tasks based on capability

Google's Pregel

- Input graph as computational structure
 - Vertex → Computing thread with local state
 - Edges → Adjacency lists stored with to the vertex
 - Communication → Push messages to (neighboring) vertices
- “Think Like a Vertex (TLV)” programming
- Easy to scale to large graphs: partition by vertex



Bulk Synchronous Programming

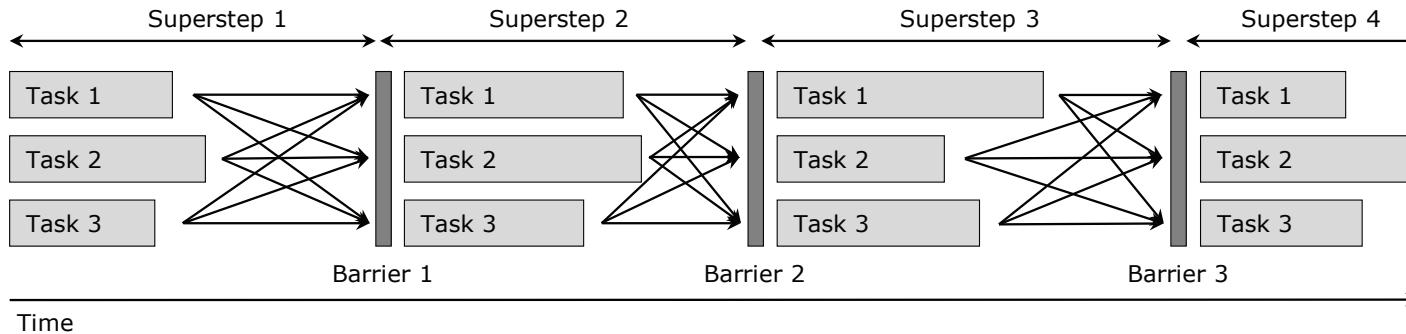


Figure taken from: [McCune et al., Thinking Like a Vertex: A Survey of Vertex-Centric Frameworks for Large-Scale Distributed Graph Processing, ACM Comp. Surveys (2015)]

Termination

- Terminate when no more active vertices
- A vertex can locally vote to **halt** and deactivate
- If halting vertex receives a message → activates

Example: S.S. Shortest Path

- **Input:** Graph (weighted edges), source vertex
- **Output:** Min distance between the source and all other vertices
- TLV implementation

vertex code:

Receive distances from neighbors, extract minimum

If minimum is smaller than current distance

 Replace current distance with minimum

For each edge

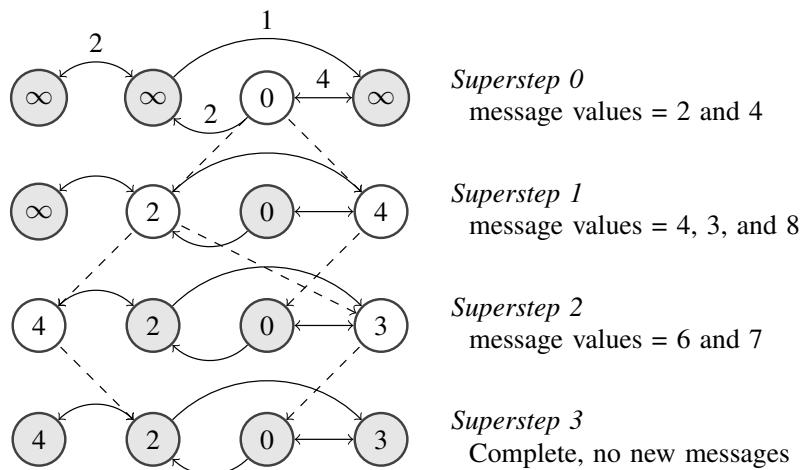
 Send current distance + edge weight

Halt

Example of SSSP Run

vertex code:

```
Receive distances from neighbors, extract minimum  
If minimum is smaller than current distance  
    Replace current distance with minimum  
For each edge  
    Send current distance + edge weight  
Halt
```



How to execute this algorithm in MapReduce?

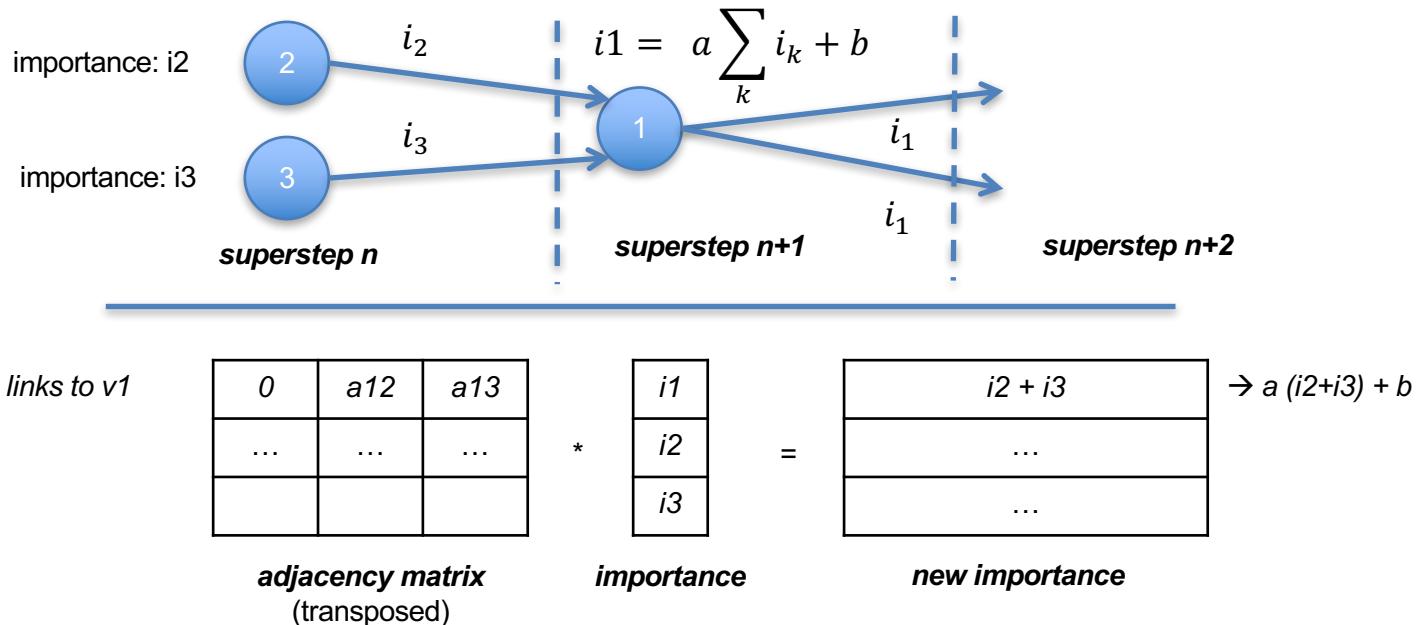
Figure taken from: [McCune et al., Thinking Like a Vertex: A Survey of Vertex-Centric Frameworks for Large-Scale Distributed Graph Processing, ACM Comp. Surveys (2015)]

Combiners and Aggregators

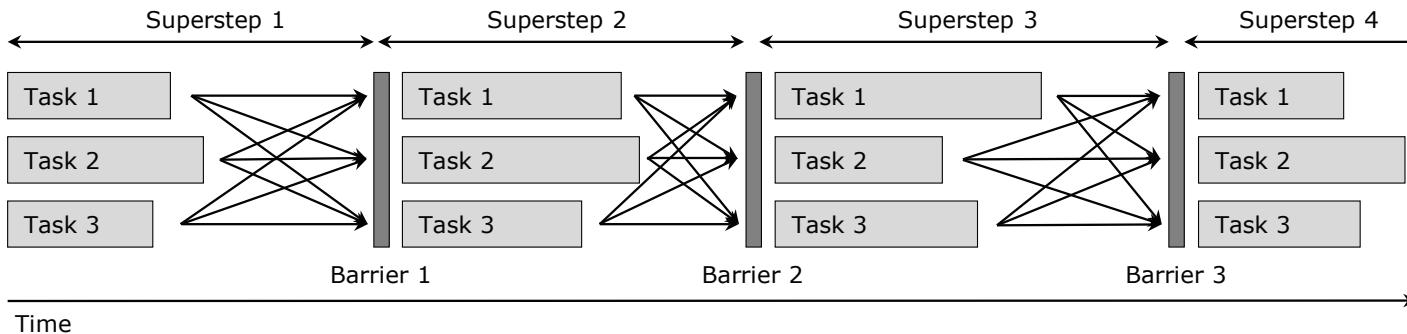
- Combiners
 - Workers merge local messages to the same vertex
 - Example in SSSP?
- Aggregators
 - Similar to map-reduce: vertices map, aggregators reduce
 - Aggregation occurs in between supersteps
 - Each worker executes local aggregation
 - Aggregation tree
 - Example: compute the distribution of distances in SSSP

Matrix-Vector Multiplication in TLV

- E.g. Page-Rank style computation



Problems with BSP?



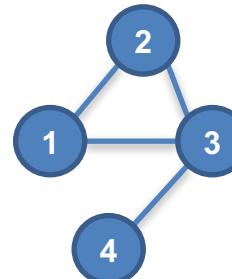
- Stragglers (workers that take longer to complete task)
- Need to balance load carefully. But how?

Asynchrony



GraphLab [2]

- Shared memory, asynchronous
- Vertices and edges have state
- Update functions
 - Modify a vertex, its incident edges, its neighbors
 - Schedule future updates on other vertices
- Serializable update functions using locking



Reference: [2] Low et al, arXiv:1006.4990, (2010)

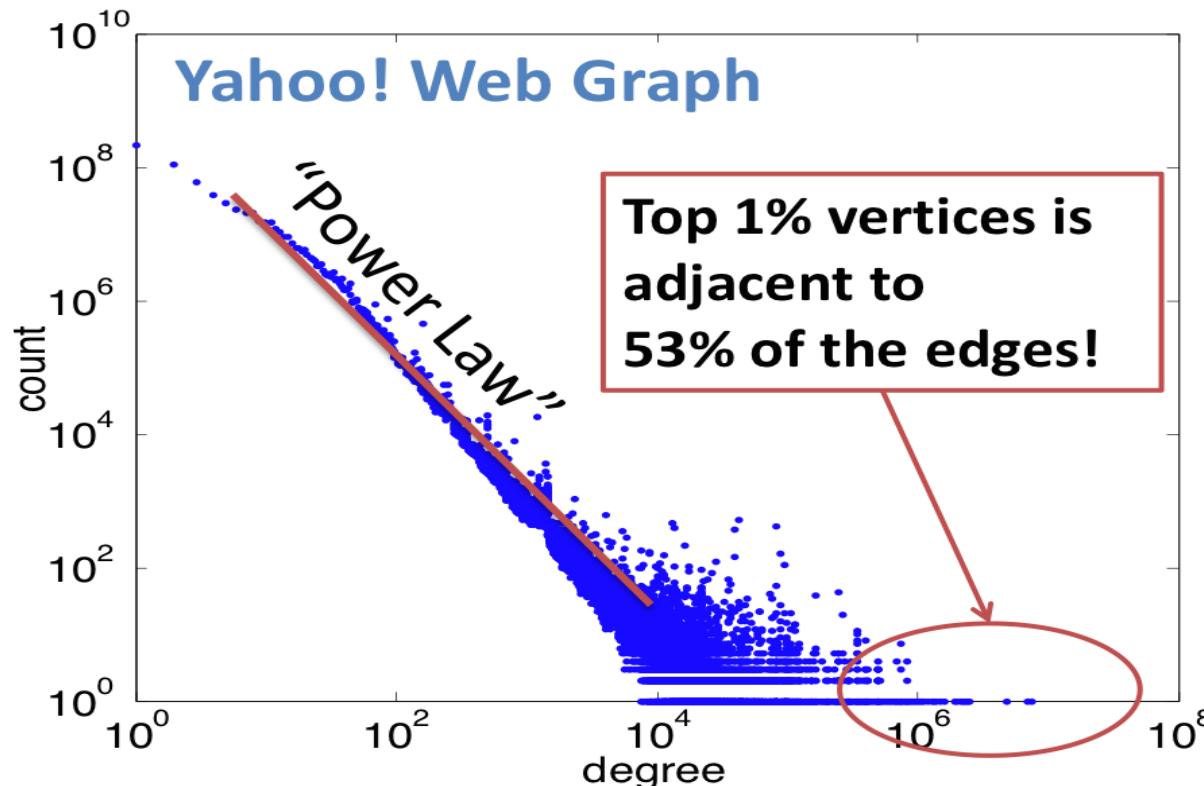
Is Asynchrony Better?

- In practice, benefits are not (yet) evident
 - Locking is expensive
 - Stragglers can still lock out other vertices
 - Application behavior is harder to understand
 - Not any longer supported even by GraphLab

HotSpots



Hotspot Vertices in Natural Graphs

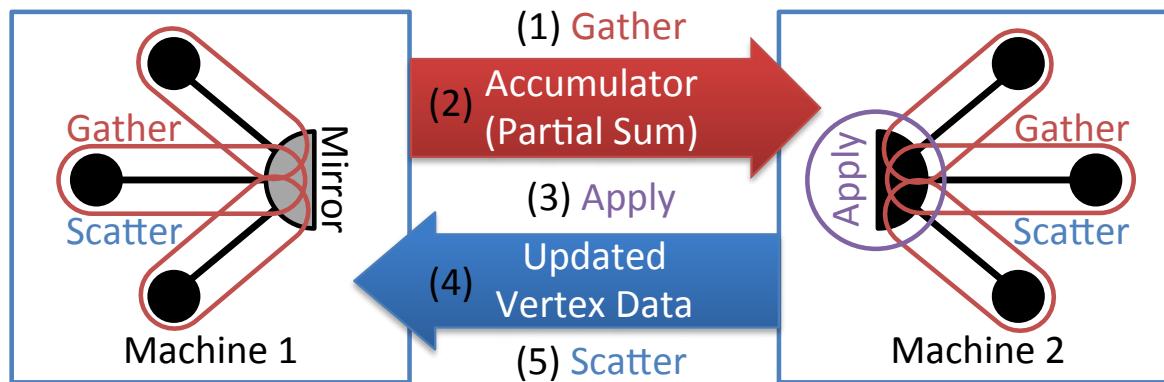


What are the issues with Hotspots?

- Long time to process all incoming messages
- Lots of output messages
- Lots of edge metadata to keep
- Delay overall execution
 - Hotspots are stragglers in synchronous systems
 - Hotspots involve lots of locking with asynchrony

PowerGraph_[3]

- Replicate high degree vertices
- Gather, Apply, Scatter



Large Graphs with 1 Server

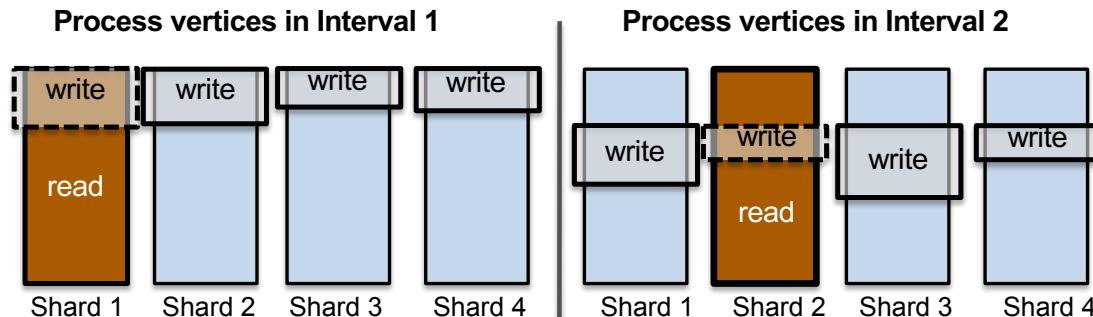


Out-Of-Core Engines

- Use a single machine
- Instead of distributing graph, store it on disk
- Efficient I/O
 - Sequential read-writes: high bandwidth and low latency
 - Random accesses: very inefficient
 - This holds not only for disks but also SSDs

GraphChi [6]

- Parallel Sliding Window
 - Sort vertices by ID and partition them in **P intervals**
 - Each interval is associated with a **shard** of edges with **destination** in the interval (sequential **reads**)
 - Sort edges in a shard by **source id** (sequential **writes**)

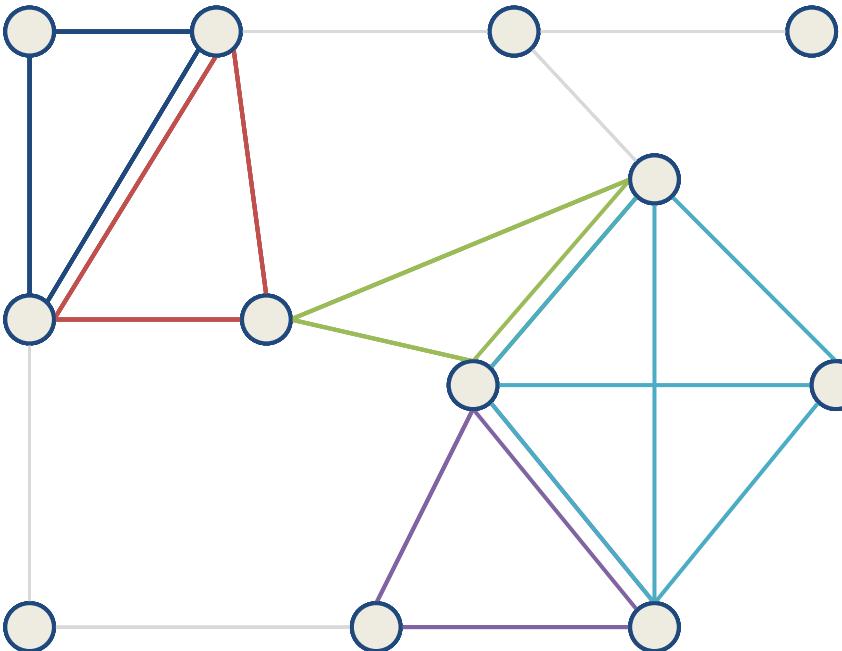


Reference: [6] Kyrola et al, OSDI 2012

Graph Mining

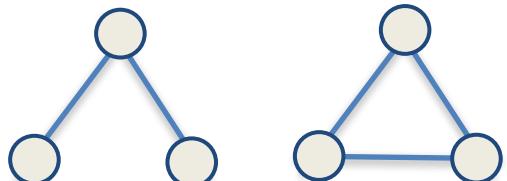


Graph Mining: Cliques

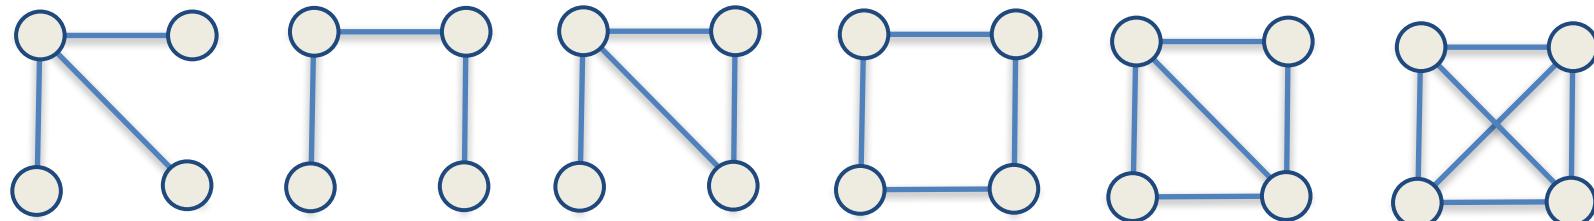


Graph Mining: Motifs

Motifs Size = 3

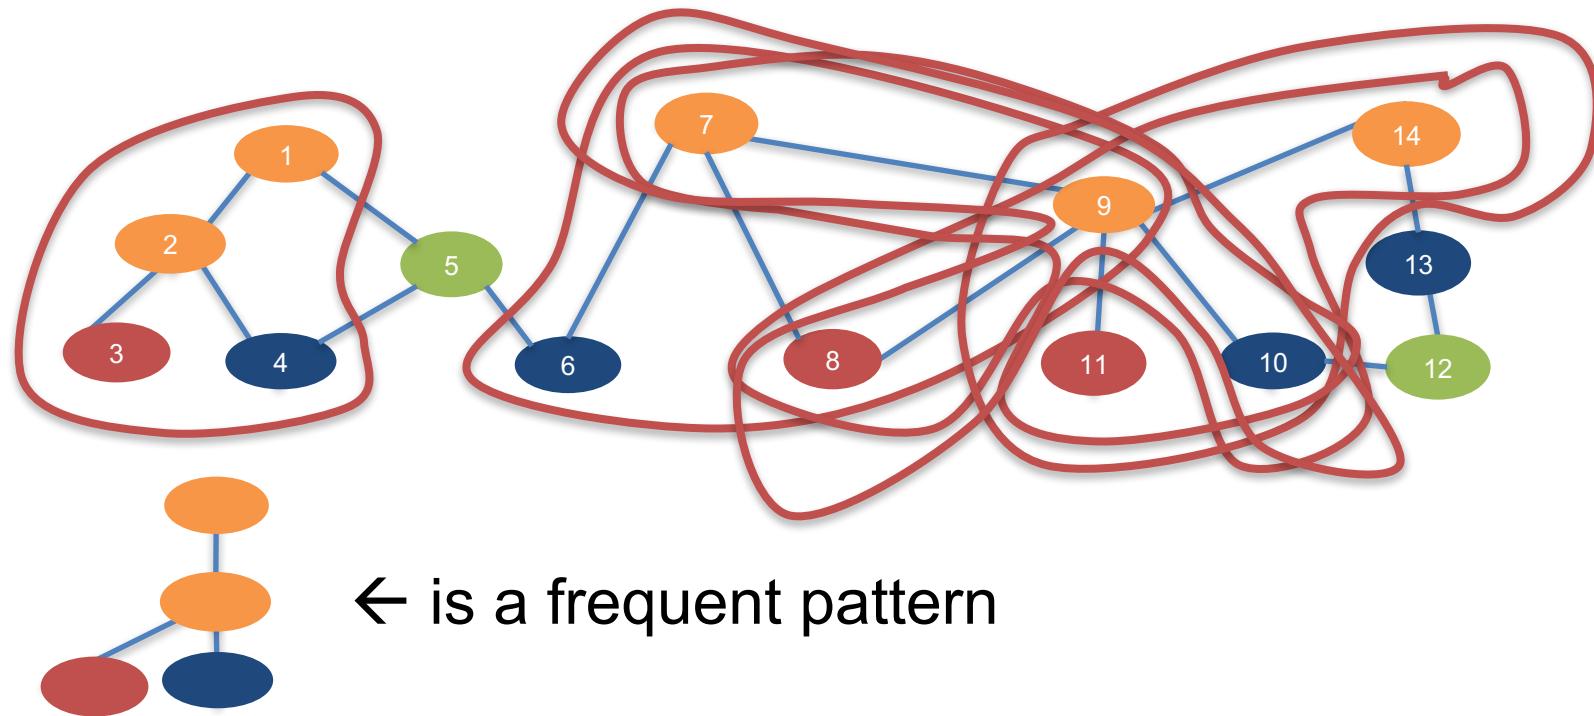


Motifs Size = 4



Graph Mining: FSM

- Frequent Subgraph mining in a single large graph.



Applications: Web and Advertising

- Link spam detection
 - Users, Web pages → Web pages
 - Dense bipartite subgraphs: Spammers
- Identify sub-markets
 - Advertisers / users → queries
 - Dense bipartite subgraphs: Advertisers / users forming sub-markets
- Attributed edges in RDF datasets
 - Infer missing edges from recurrent patterns



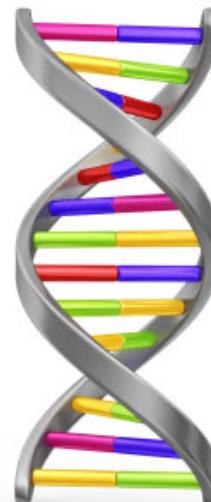
Applications: Social Computing

- Friend recommendation
 - Missing edges in quasi-cliques
- Community detection
- Finding stories in social media
 - Vertices: Entities (e.g. politicians)
 - Edges: Co-occurrence in tweets
 - Frequently co-occurring entities are involved in a story

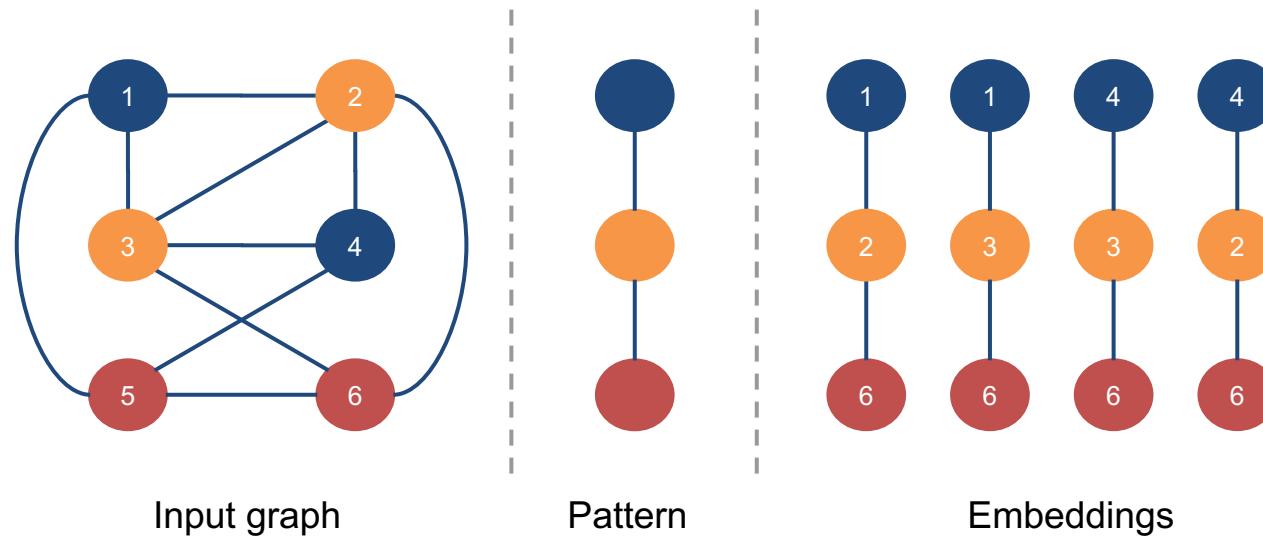


Applications: Biology

- DNA motif detection
- Gene correlation
- Protein-protein interaction

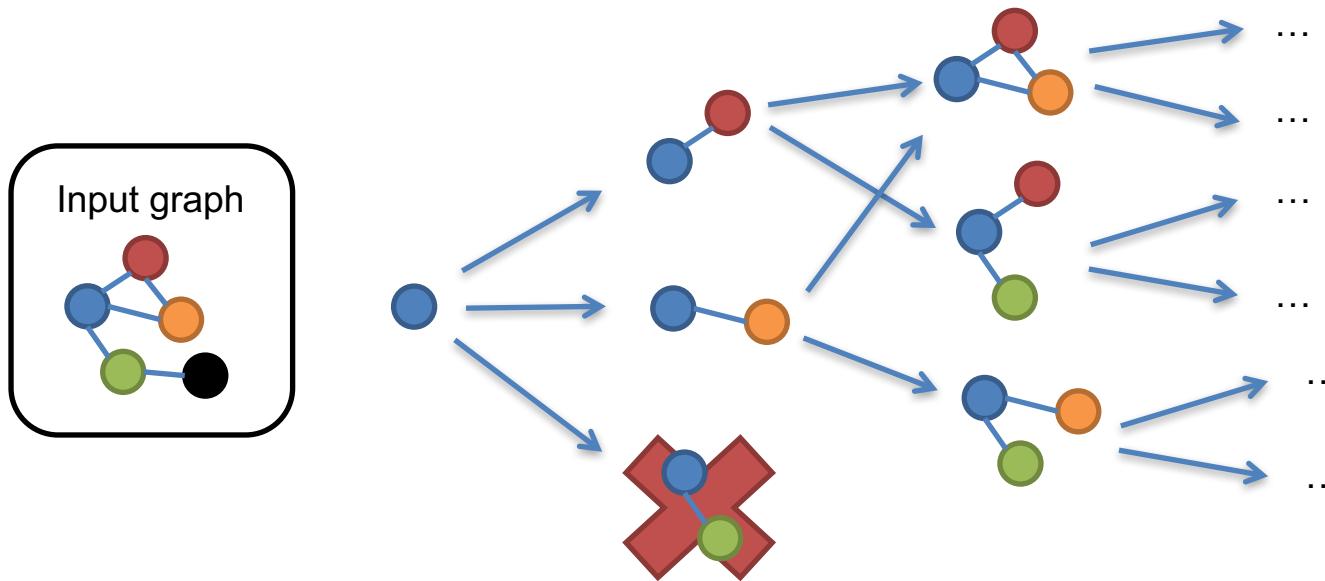


Graph Mining - Concepts

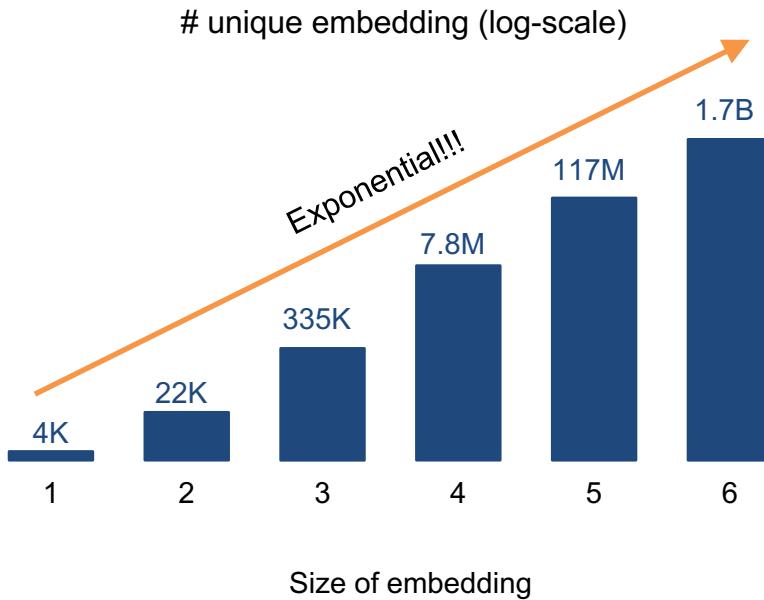
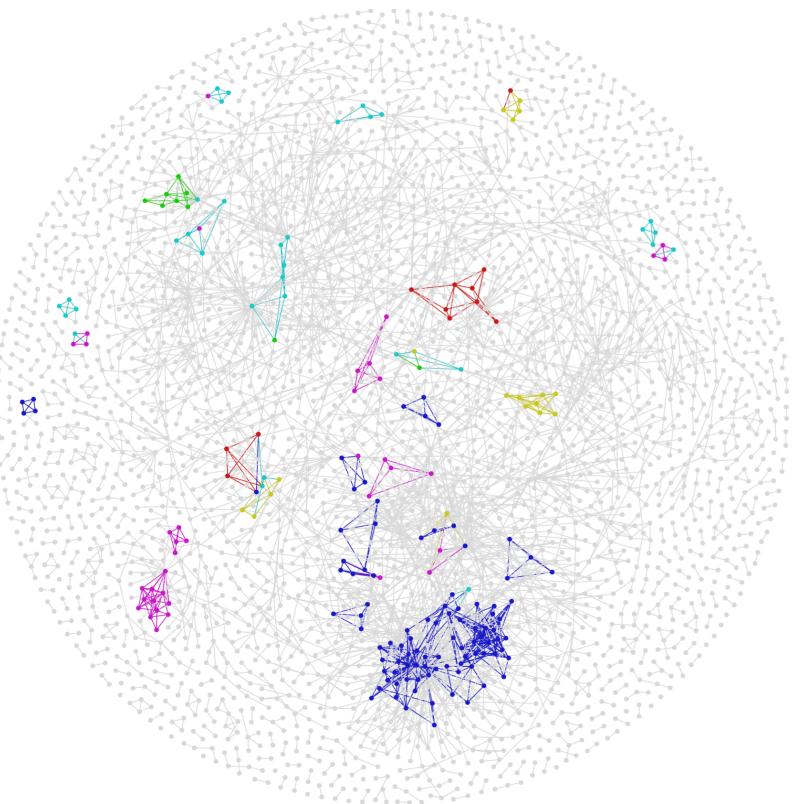


Graph Exploration

- Enumerate (& prune) embeddings
- Aggregate by pattern



Challenges



Exponential number of embeddings

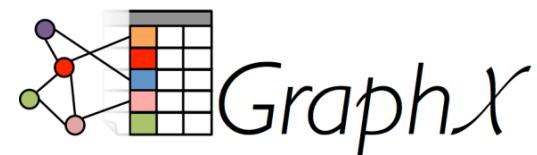
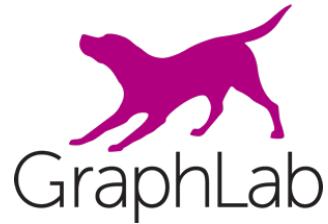
State of the Art: Custom Algorithms

	Easy to Code	Efficient Implementation	Transparent Distribution
Custom Algorithms	X	✓	X
	Easy to Code	Efficient Implementation	Transparent Distribution



State of the Art: Think Like a Vertex

	Easy to Code	Efficient Implementation	Transparent Distribution
Custom Algorithms	✗	✓	✗
Think Like a Vertex	✗	✗	✓



Arabesque [7]

- New execution model & system
 - Think Like an Embedding
 - Purpose-built for distributed graph mining
 - Hadoop-based
- Contributions:
 - Simple & Generic API
 - High performance
 - Distributed & Scalable by design



Arabesque

	Easy to Code	Efficient Implementation	Transparent Distribution
Custom Algorithms	✗	✓	✗
Think Like a Vertex	✗	✗	✓
Arabesque	✓	✓	✓



API Example: Clique finding

```
1 boolean filter(Embedding e) {  
2     return isClique(e);  
3 }  
4 void process(Embedding e) {  
5     output(e);  
6 }  
7 boolean shouldExpand(Embedding embedding) {  
8     return embedding.getNumVertices() < maxsize;  
9 }  
10 boolean isClique(Embedding e) {  
11     return e.getNumEdgesAddedWithExpansion()==e.getNumberOfVertices()-1;  
12 }
```

State of the Art
(Mace, centralized)

4,621 LOC

API Example: Motif Counting

```
1 boolean filter(Embedding e) {  
2     return true;  
3 }  
4 void process(Embedding embedding) {  
5     output(embedding);  
6     map(AGG_MOTIFS, embedding.getPattern(), reusableLongWritableUnit);  
7 }  
8 boolean shouldExpand(Embedding embedding) {  
9     return embedding.getNumVertices() < maxsize;  
10 }
```

State of the Art
(GTrieScanner, centralized)

3,145 LOC

API Example: FSM

- Ours was the first distributed implementation
- 280 lines of Java Code
 - ... of which 212 compute frequent metric
- Baseline (GRAMI): 5,443 lines of Java code.

Arabesque: An Efficient System

- As efficient as centralized state of the art

Application - Graph	Centralized Baseline	Arabesque 1 thread
Motifs - MiCo (MS=3)	50s	37s
Cliques - MiCo (MS=4)	281s	77s
FSM - CiteSeer (S=300)	4.8s	5s

Arabesque: A Scalable System

- Scalable to thousands of workers
- Hours/days → Minutes

Application - Graph	Centralized Baseline	Arabesque 640 cores
Motifs	2 hours 24 minutes	25 seconds
Cliques	4 hours 8 minutes	1 minute 10 seconds
FSM - Patents	> 1 day	1 minute 28 seconds

First Distributed Implementation

How: Arabesque Optimizations

- Avoid Redundant Work
 - Efficient canonicity checking
- Subgraph Compression
 - Overapproximating Directed Acyclic Graphs (ODAGs)
- Efficient Aggregation
 - 2-level pattern aggregation

Arabesque fundamentals

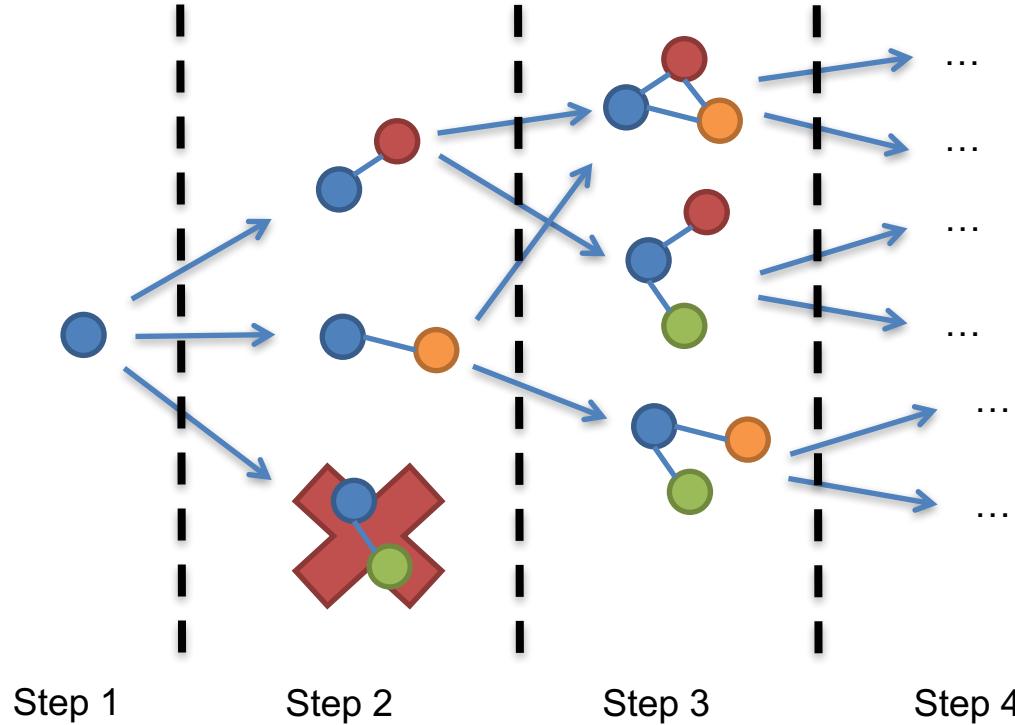
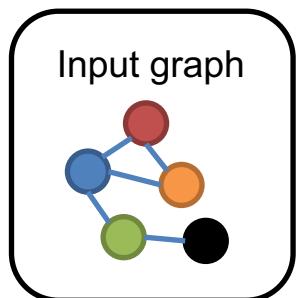


Arabesque: Fundamentals

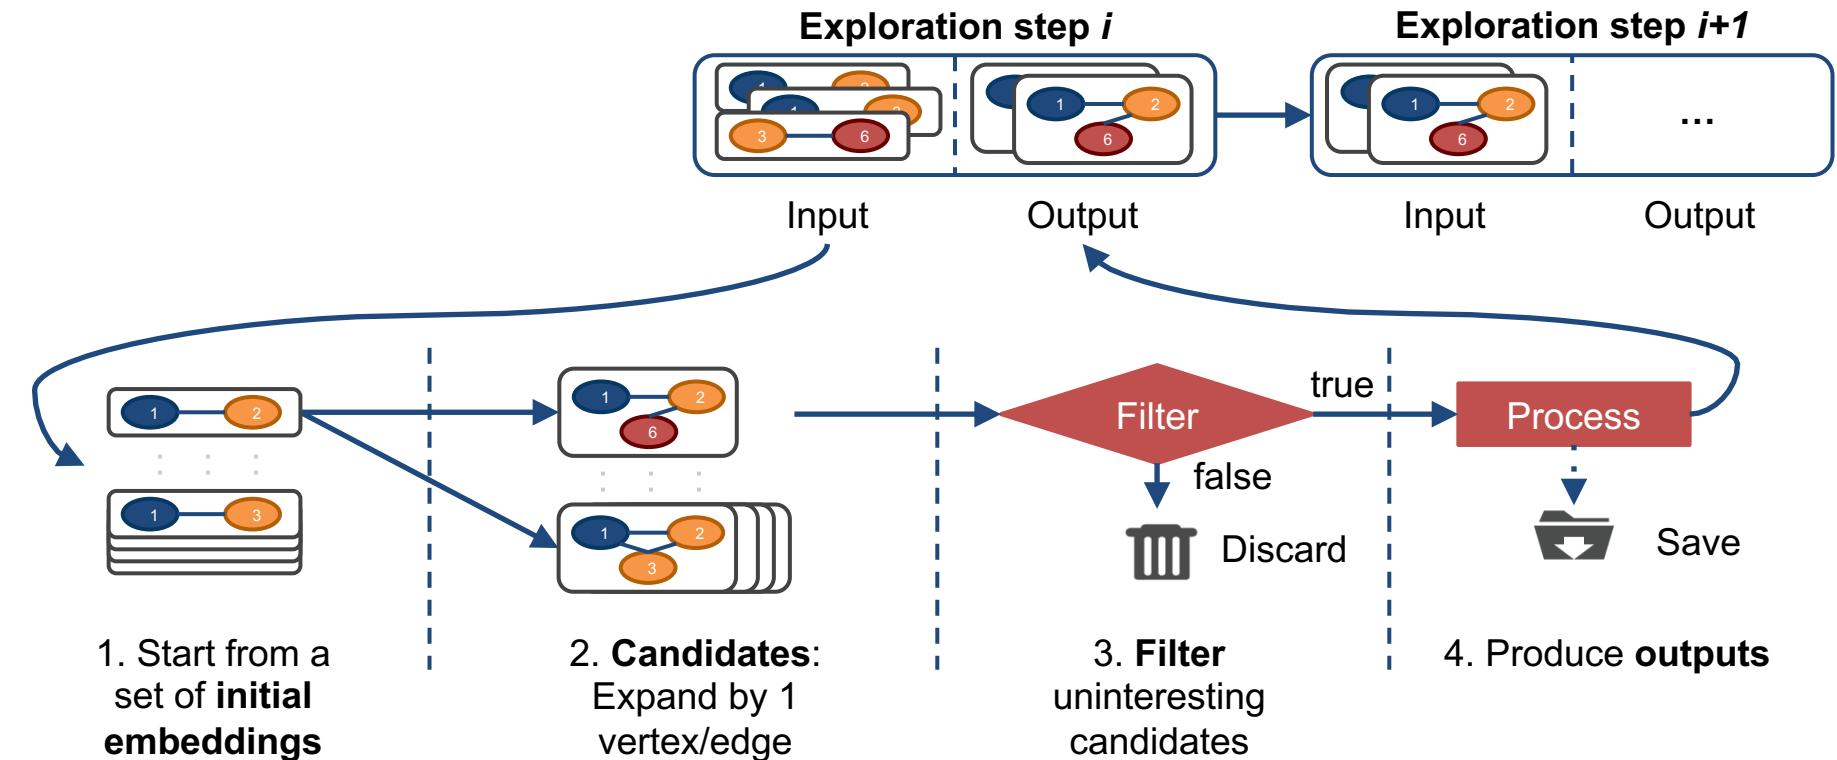
- Embeddings as 1st class citizens:
 - **Think Like an Embedding** model
 - Functional programming

Arabesque responsibilities	User responsibilities: define
<div>Graph Exploration</div> <div>Load Balancing</div>	<div><i>filter</i> (Embedding e)</div> <div><i>process</i> (Embedding e)</div>

Graph Exploration



Model - Think Like an Embedding



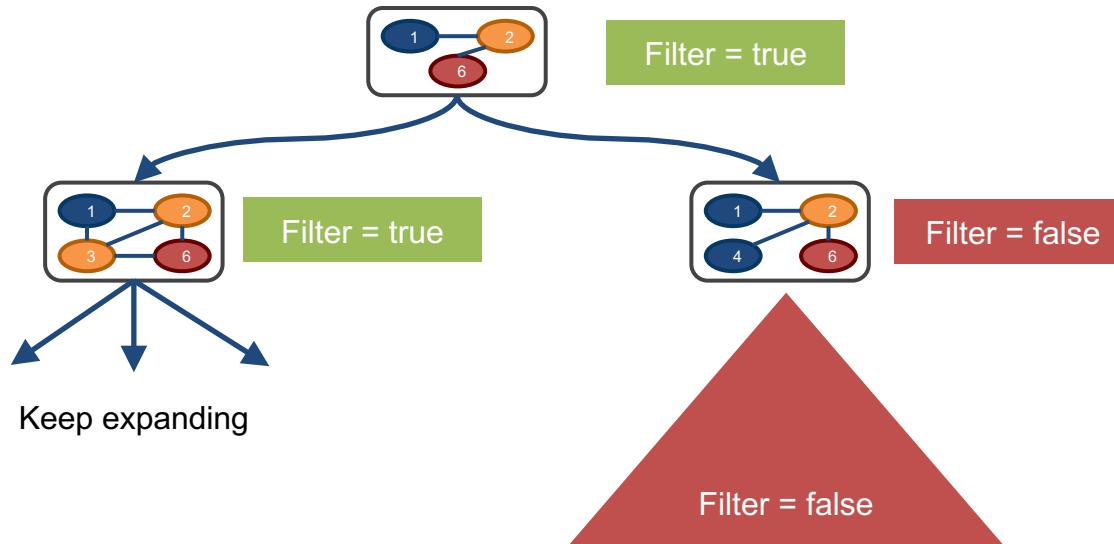
API Example: Clique finding

```
1 boolean filter(Embedding e) {
2     return isClique(e);
3 }
4 void process(Embedding e) {
5     output(e);
6 }
7 boolean shouldExpand(Embedding embedding) {
8     return embedding.getNumVertices() < maxsize;
9 }
10 boolean isClique(Embedding e) {
11     return e.getNumEdgesAddedWithExpansion()==e.getNumberOfVertices()-1;
12 }
```

Guarantee: Completeness

For each e , if $\text{filter}(e) == \text{true}$ then $\text{Process}(e)$ is executed

Requirement: Anti-monotonicity



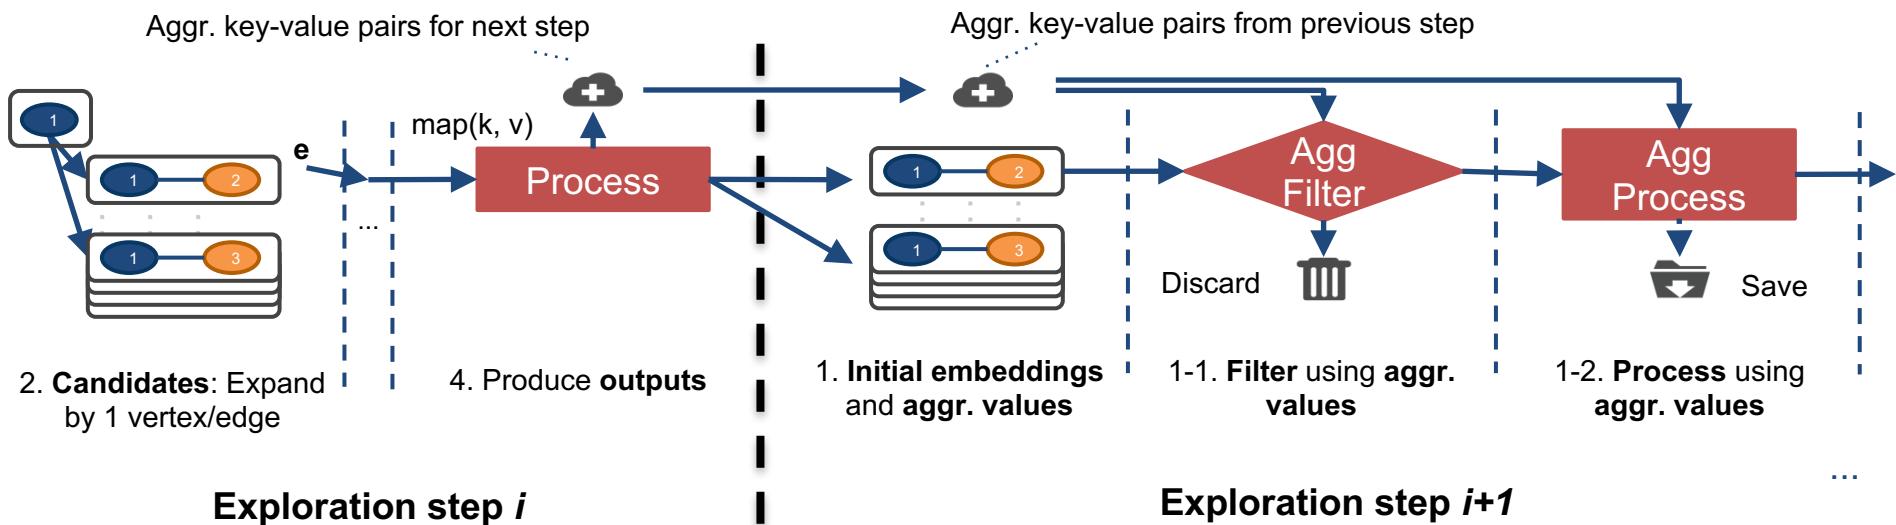
We can **prune** and be sure that we won't ignore
desired embeddings

Aggregation during expansion

- Filter might need aggregated values
 - E.g.: Frequent subgraph mining
 - Frequency calculation → look at all candidates
- Aggregation in parallel with exploration step
 - Embeddings filtered as soon as aggregated values are ready.

Aggregation during expansion

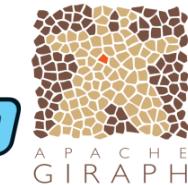
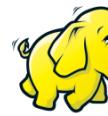
- Filter function may depend on aggregated data across all embeddings of the same pattern.



System Architecture & Optimizations



Arabesque Architecture



Previous step

Input
Embeddings
size n

split 1
split 4
split 7



Worker 1



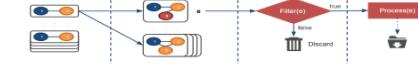
Output
Embeddings size
 $n + 1$

split 1
split 4
split 7

split 2
split 5
split 8



Worker 2



split 2
split 5
split 8

split 3
split 6
split 9



Worker 3

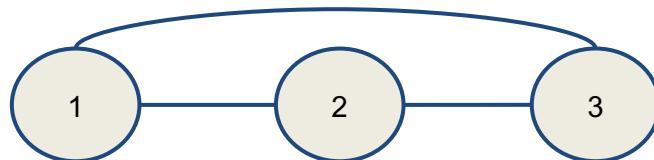


split 3
split 6
split 9

Next step

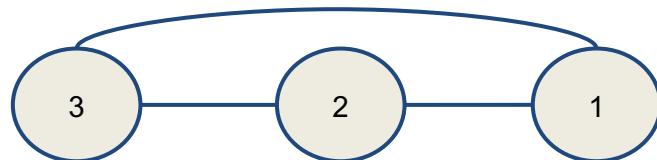
Avoiding redundant work

- **Problem:** Automorphic embeddings
 - Automorphisms == subgraph equivalences
 - Redundant work



Worker 1

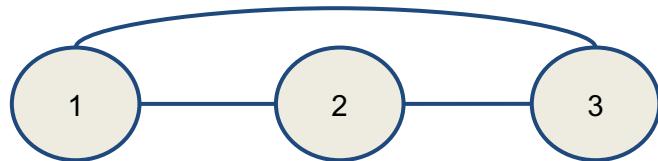
=



Worker 2

Avoiding redundant work

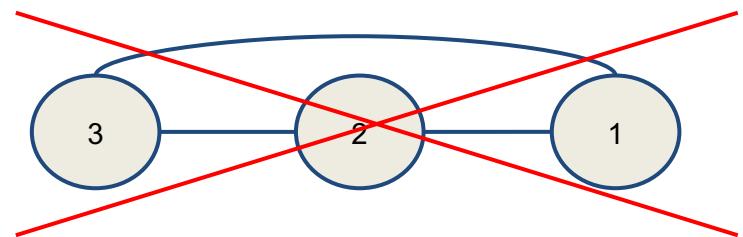
- **Solution: Decentralized Embedding Canonicity**
 - No coordination
 - Efficient



Worker 1

`isCanonical(e) → true`

\equiv

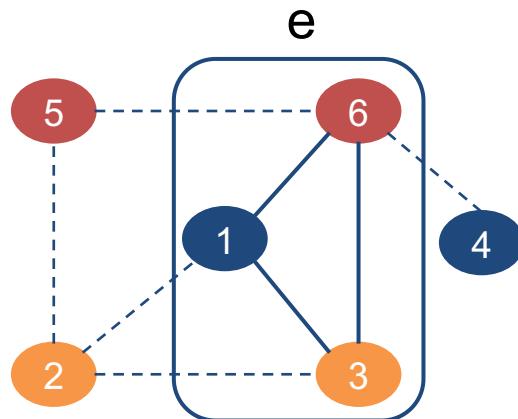


Worker 2

`isCanonical(e) → false`

Embedding Canonicality

- $\text{isCanonical}(e)$ iff at every step add neighbor with smallest ID



Initial embedding (e)

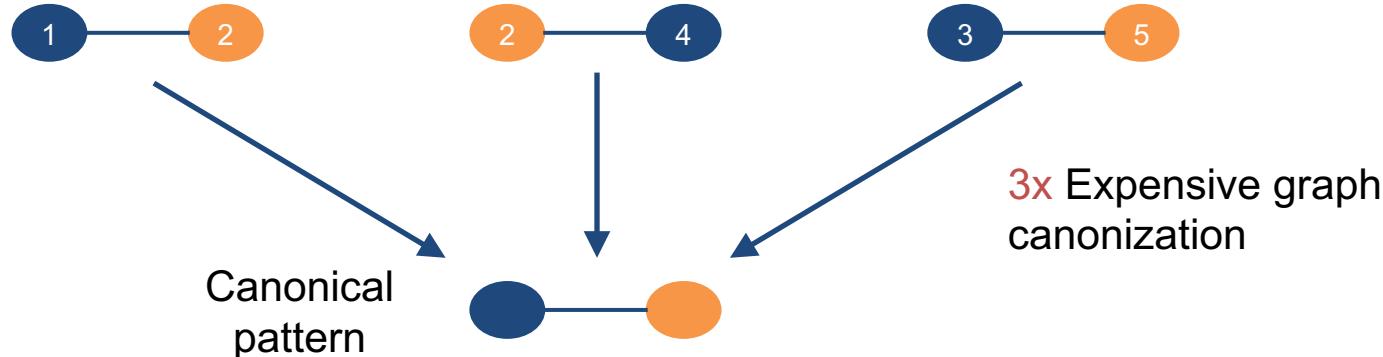
- 1 - 3 - 6

Expansions:

- 1 - 3 - 6 - 5 → canonical
- 1 - 3 - 6 - 4 → canonical
- 1 - 3 - 6 - 2 → not canonical (1 - 2 - 3 - 6)

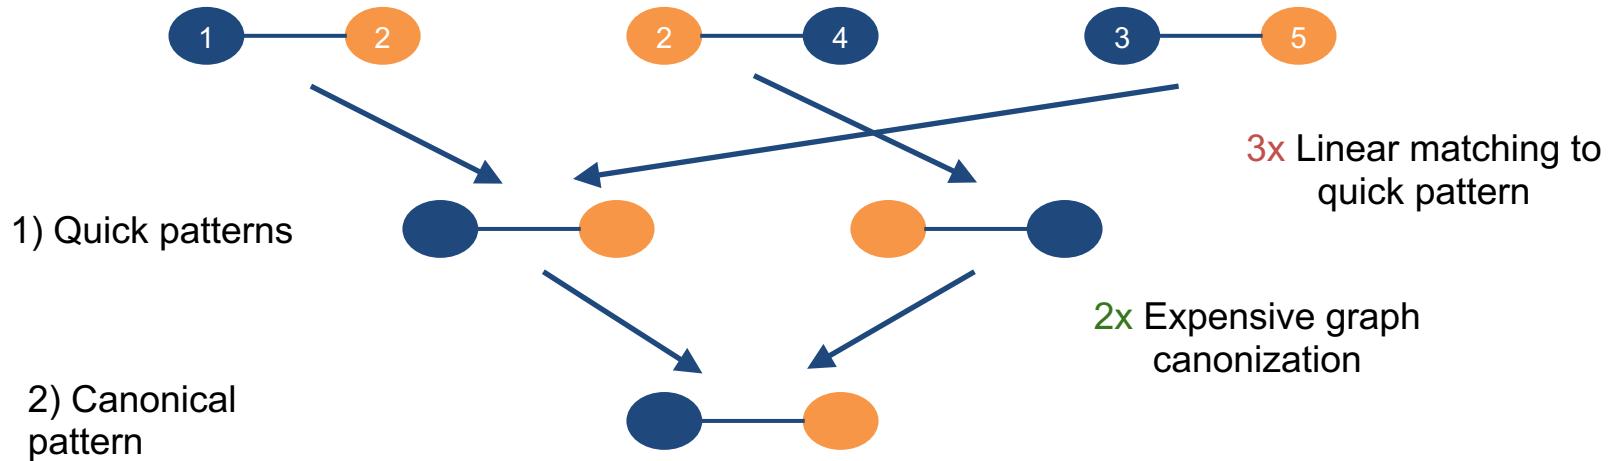
Efficient Pattern Aggregation

- **Goal:** Aggregate automorphic patterns to single key
 - Find canonical pattern
 - No known polynomial solution



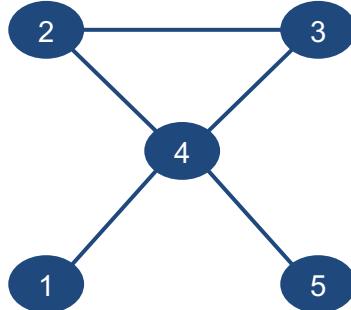
Efficient Pattern Aggregation

- **Solution:** 2-level pattern aggregation
 1. Embeddings → quick patterns
 2. Quick patterns → canonical pattern



Handling Exponential growth

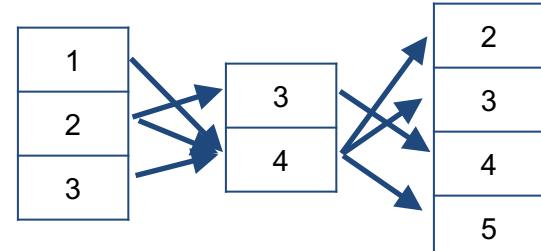
- **Goal:** handle trillions+ different embeddings?
- **Solution: Overapproximating DAGs (ODAGs)**
 - Compress into less restrictive superset
 - Deal with spurious embeddings



Input Graph

Canonical Embeddings		
1	4	2
1	4	3
1	4	5
2	3	4
2	4	5
3	4	5

Embedding List



ODAG

Evaluation



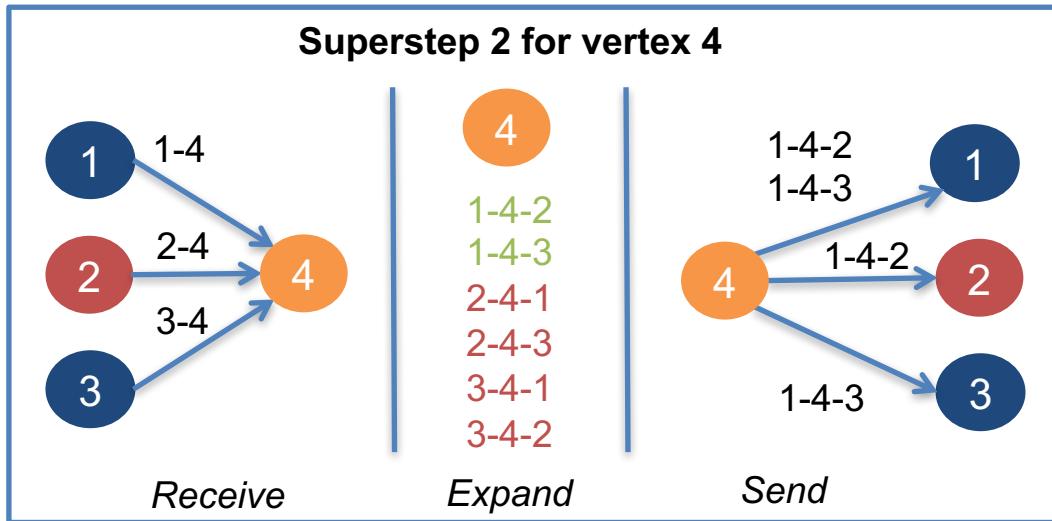
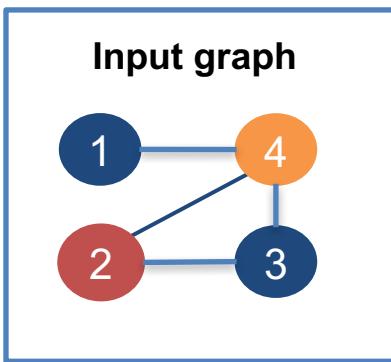
Evaluation - Setup

- 20 servers: 32 threads @ 2.67 GHz, 256GB RAM
- 10 Gbps network
- 3 algorithms: Frequent Subgraph Mining, Counting Motifs and Clique Finding
- Input graphs:

	# Vertices	# Edges	# Labels	Avg. Degree
CiteSeer	3,312	4,732	6	3
MiCO	100,000	1,080,298	29	22
Patents	2,745,761	13,965,409	37	10
Youtube	4,589,876	43,968,798	80	19
SN	5,022,893	198,613,776	0	79
Instagram	179,527,876	887,390,802	0	10

Graph Exploration with TLV

1. Receive embeddings
2. Expand by adding neighboring vertices
3. Send *canonical* embeddings to their constituting vertices

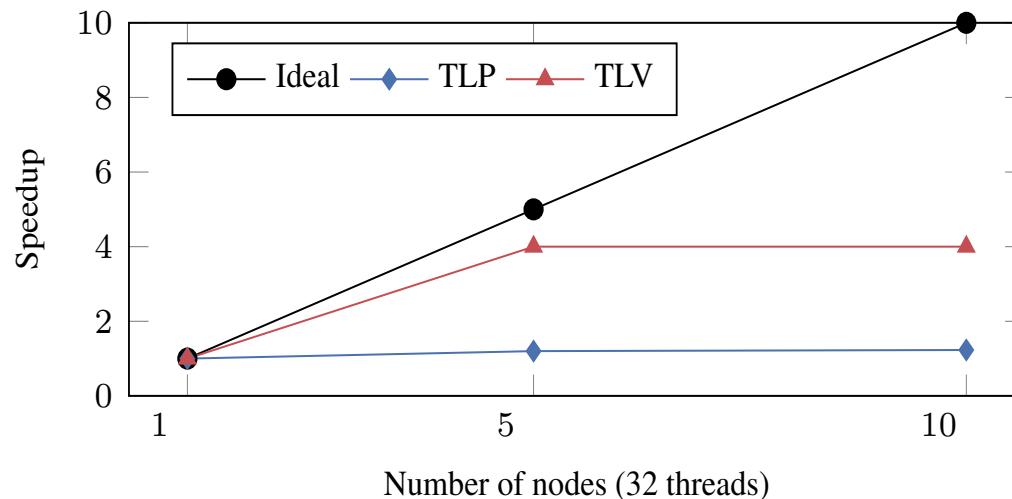


Think Like a Pattern

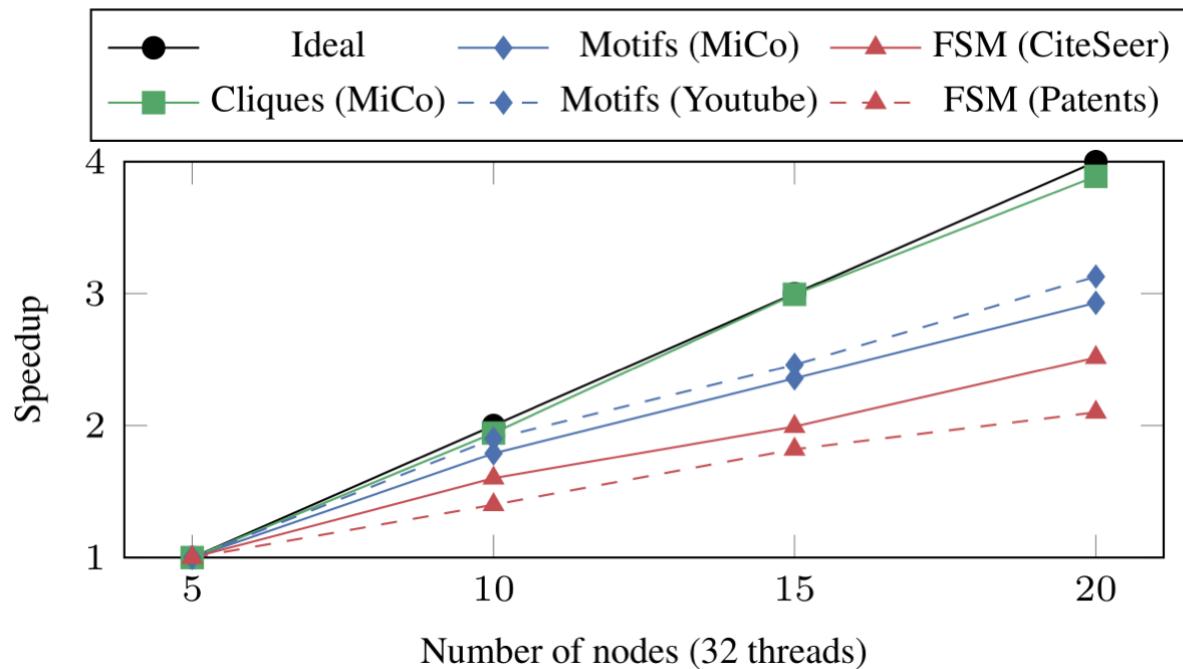
- Many existing algorithms keep state by pattern
- Advantages
 - Rebuild embeddings from scratch
 - No need to materialize full intermediate state
- Idea of TLP:
 - Assign different patterns to different machines
 - Avoid storing materialized embedding

Evaluation - TLP & TLV

- Use case: frequent subgraph mining
- No scalability. Bottlenecks:
 - TLV: Replication of embeddings, hotspots
 - TLP: very few patterns do all the work



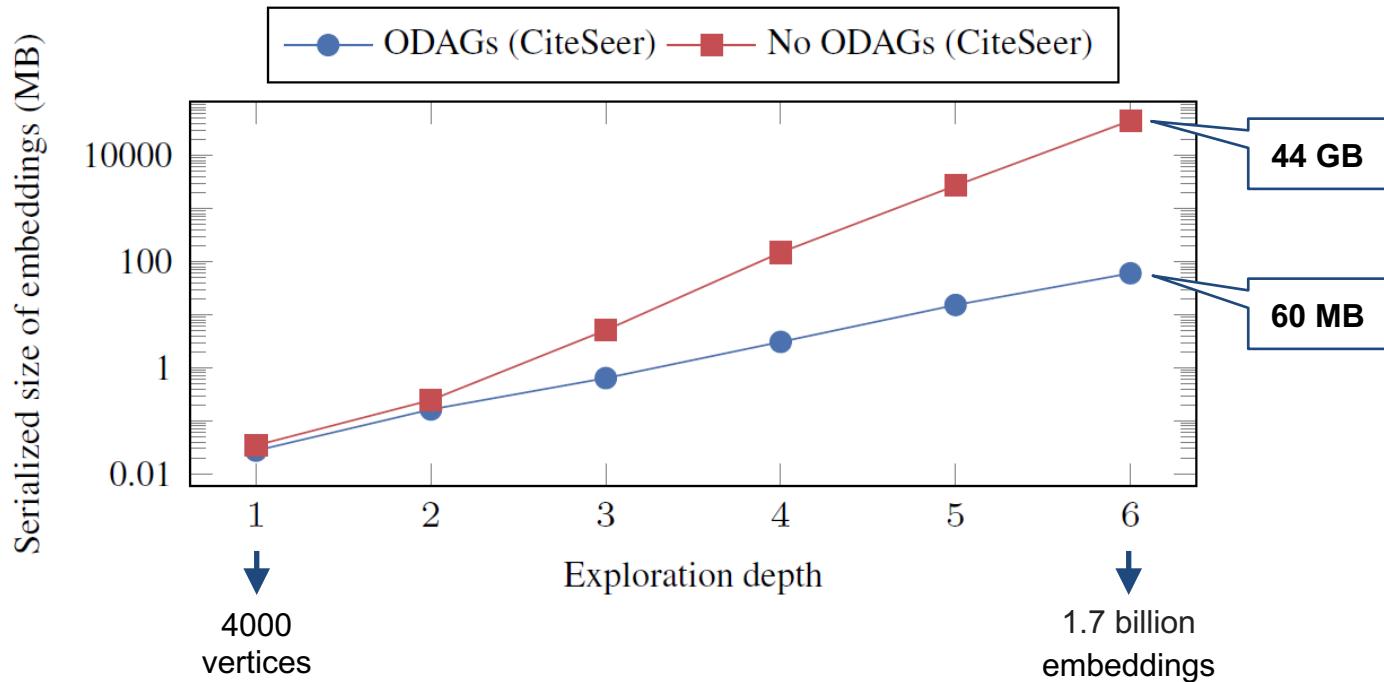
Evaluation – Arabesque Scalability



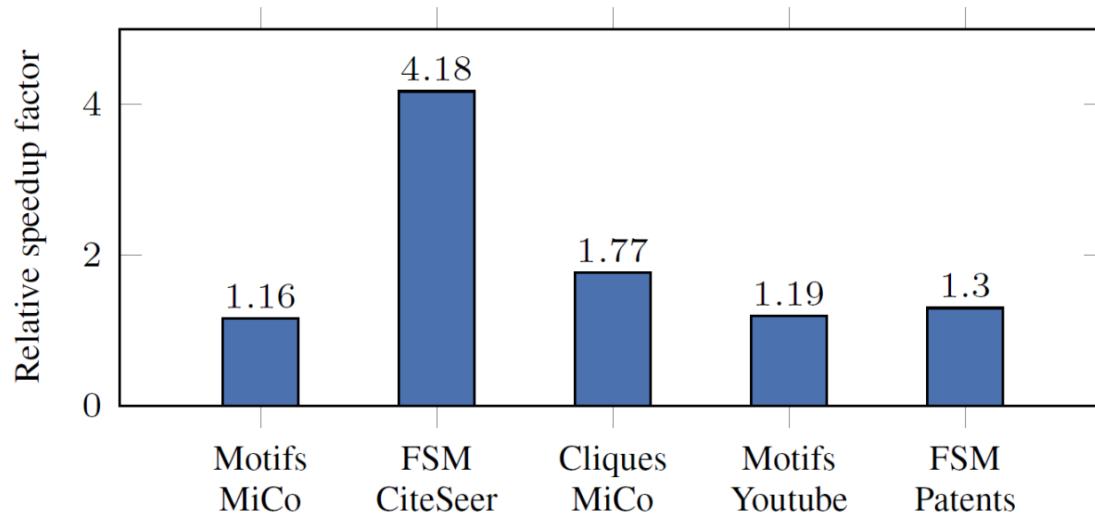
Evaluation – Arabesque Scalability

Application - Graph	Centralized Baseline	Arabesque - Num. Servers (32 threads)				
		1	5	10	15	20
Motifs - MiCo	8,664s	328s	74s	41s	31s	25s
FSM - Citeseer	1,813s	431s	105s	65s	52s	41s
Cliques - MiCo	14,901s	1,185s	272s	140s	91s	70s
Motifs - Youtube	Fail	8,995s	2,218s	1,167s	900s	709s
FSM - Patents	>19h	548s	186s	132s	102s	88s

Evaluation - ODAGs Compression

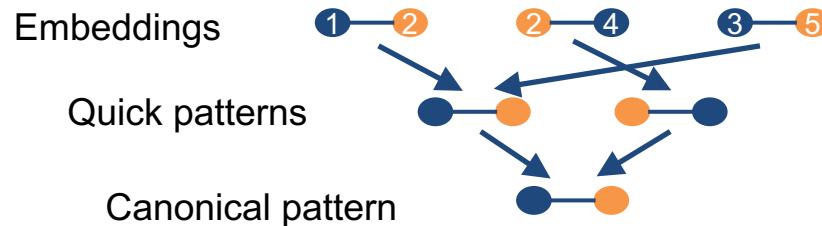


Evaluation - Speedup w ODAGs

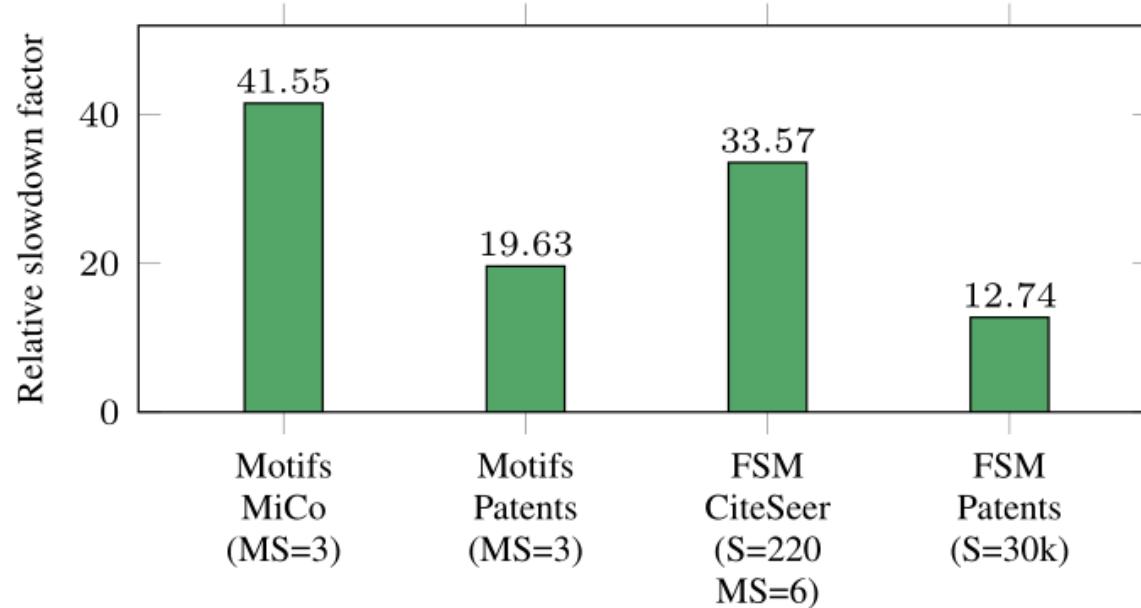


Evaluation - 2-level aggregation

	Motifs MiCo (MS = 4)	Motifs Youtube (MS=4)	FSM CiteSeer (S=220, MS=7)	FSM Patents (S=24k)
Embeddings	10,957,439,024	218,909,854,429	1,680,983,703	1,910,611,704
Quick Patterns	21	21	1433	1800
Canonical Patterns	6	6	97	1348
Reduction Factor	521,782,810x	10,424,278,782x	1,173,052x	1,061,451x



Evaluation - 2-level aggregation



What's Next?



Download It, Play with It, Hack It



Distributed Graph Mining
Made Easy
[Documentation](#)
[How to Use](#)

Distributed Graph Mining Made Easy

Arabesque is a distributed graph mining system that enables quick and easy development of graph mining algorithms, while providing a scalable and efficient implementation that runs on top of Hadoop.

Benefits of Arabesque:

- Simple intuitive API, tuned for Graph Mining Problems
- Handles all the complexity of Graph Mining Algorithms transparently
- Scalable to hundreds of machines
- Efficient implementation: negligible overhead compared to equivalent centralized solutions
- Support of large graphs with over a billion edges. It can process trillion of subgraphs in a commodity cluster.
- Designed for Hadoop. Runs as an Apache Giraph Job.
- Open-Source with Apache 2.0 license.

Documentation

Check our SOSP 2015 [paper](#) that describe the system.

Follow our [user-guide](#), on how to program graph mining applications on Arabesque.

How to Use

Binary jars can be downloaded [here](#).

The source code can be accessed from [github](#).

<http://arabesque.io>

- **Open-source (Apache 2.0)**
- Pre-compiled jar
- User guide

Future Directions

- Full support for graph search + mining
 - TLE simplifies the design of distributed graph search
- Support for real-time graph data
 - Incremental computations
- Multiplatform
 - Spark – ongoing
 - Flink - planned
- Library of graph mining applications

How to Run & Code



Requirements

- Hadoop installation:
 - Runs a map-reduce job (Giraph based)
- To develop:
 - Java 7

Input Graph

- Graphs:
 - labels on vertices
 - labels on edges
 - Multiple edges with labels between two vertices
- Graph should have sequential vertex ids, and it should be ordered

How to Run?

```
./run_arabesque.sh cluster.yaml application.yaml
```

API Example: Clique finding

```
1 boolean filter(Embedding e) {  
2     return isClique(e);  
3 }  
4 void process(Embedding e) {  
5     output(e);  
6 }  
7 boolean shouldExpand(Embedding embedding) {  
8     return embedding.getNumVertices() < maxsize;  
9 }  
10 boolean isClique(Embedding e) {  
11     return e.getNumEdgesAddedWithExpansion()==e.getNumberOfVertices()-1;  
12 }
```

State of the Art
(Mace, centralized)

4,621 LOC

Cluster.yaml

```
num_workers: 10
num_compute_threads: 16
output_active: yes

# Giraph configuration
#giraph.nettyClientThreads: 32
#giraph.nettyServerThreads: 32
#giraph.nettyClientExecutionThreads: 32
#giraph.channelsPerServer: 4
#giraph.useBigDataIOForMessages: true
#giraph.useNettyPooledAllocator: true
#giraph.useNettyDirectMemory: true
#giraph.nettyRequestEncoderBufferSize: 1048576
```

Cliques.yaml

```
computation: io.arabesque.examples.clique.CliqueComputation
input_graph_path: citeseer-single-label.graph
output_path: Cliques_Output

#communication_strategy: embeddings

# Custom parameters
arabesque.clique.maxsize: 4
```